

1 Recap

Abstract Data Type (ADT) an object with well-defined operations, e.g. a stack supports push() and pop() operations

stack can be implemented using different data structures, e.g. a stack can be implemented using a linked list or an array

Array a contiguous sequence of objects with the same size close to how computers store data in their memory can also be multi-dimensional

support fast access to their elements through indexing inserting and inserting values in arbitrary locations are expensive

in python: no built-in data structure, lists are indexable, use numpy library for proper/faster arrays

Lists main operations: append/prepend, head/tail

typically implemented using linked lists (python lists are array-based)

Stacks last-in-first-out (LIFO) data structure

basic operations: push, pop

can be implemented using linked lists (or arrays)

Queues first-in-first-out (FIFO) data structure

basic operations: enqueue, dequeue

can be implemented using linked lists (or maybe arrays)

Strings often implemented based on character arrays

Maps/Dictionarys similar to arrays and lists, but allow indexing with (almost) arbitrary data types

maps generally implemented using hashing

Sets implement the mathematical (finite) sets, a collection of unique elements without order

Algorithms Desirable properties: correctness, robustness, efficiency, simplicity

Resursion have to define one or more base cases

each recursive step should approach the base case

2 Analysis of Algorithms

characterize running times of algorithms as a function of input size:

Family	Definition
Constant	f(n) = c
Logarithmic	f(n) = log _b n
Linear	f(n) = n
N log N	f(n) = n log n
Quadratic	f(n) = n ²
Cubic	f(n) = n ³
Other polynomials	f(n) = n ^k , for k > 3
Exponential	f(n) = b ⁿ , for b > 1
Factorial	f(n) = n!

Logarithm the inverse of exponentiation: x = log_b n ->

b^{log_b n = n}

for us, no base means base-2

log xy = log x + log y

log (x/y) = log x - log y

log x^a = a log x

log_b x = log_a x / log_b a

grow much slower than linear functions

Polynomials degree-0: f(n) = c

degree-1: f(n)=n+c

degree-2: f(n)=n²+n+c

generally drop the lower order terms

1+2+3+...+n = n(n+1)/2

Permutations

n! = n*(n-1)*...*2*1

P(n,k) = n*(n-1)*...*(n-k-1) = n!/(n-k)!

Combinations

n choose k C(n,k) = P(n,k)/P(k,k) = n!/(n-k)!k!

P(n,k) = n*(n-1)*...*(n-k-1) = n!/(n-k)!

Proof by induction

used for both proving the correctness and running times of algorithms

show that the base case holds

assume the result is correct for n, show that it also holds for n+1

Hardware independence characterized by random access memory (RAM)

the data and instructions are stored in the RAM

processing unit performs basic operations in constant time any memory cell with an address can be accessed in equal (constant) time

the processor fetches them as needed and executes following the instruction

there may be other, specialized registers

modern processing units also employ a cache

Formal analysis of running time simply count

the number of primitive operations

primitive operations include: assignment, arithmetic operations, compare primitive data types (in Java: boolean, byte, char, short, int, long, float and double), access a single

memory location, function calls, return from functions non-primitive operations: loops, recursion, compare sequences

Big-O notation used for indicating an upper bound of an algorithm as a function of running time

if running time of an algorithm is O(f(n)), its running time grows proportional to f(n) as the input size n grows

drop the constants and lower order terms

transitivity: if f(n)=O(g(n)), and g(n)=O(h(n)), then f(n)=O(h(n))

additivity: if both f(n) and g(n) are O(h(n)), f(n)+g(n)=O(h(n))

maximum problem size asymptotic analysis is important: assume we can solve a problem of size m in a given time on current hardware, when we get a better computer, we can calculate the new problem size we can solve in the same time

gap between polynomial and exponential algorithms: problem size for exponential algorithms does not scale with faster computers

worst case analysis

in most cases we are interested in the worst case analysis average case analysis is also useful, but requires defining a distribution over possible inputs and is often more challenging

pro: easier, get a very strong guarantee that the algorithm won't perform worse than the bound

con: in some problems, worst case examples are very rare

asymptotic analysis

our analyses are based on asymptotic behavior

pro: correct for a 'large enough' input

con: constant or lower order factors are not always unimportant

Big-O relatives Big-O(upper bound): f(n) is O(g(n)) if f(n) is asymptotically less than or equal to g(n)

Big-Omega (lower bound): f(n) is Omega(g(n)) if f(n) is asymptotically greater than or equal to g(n)

Big-Theta (upper/lower bound): f(n) is Theta(g(n)) if f(n) is asymptotically equal to g(n): f(n) is O(g(n)) and f(n) is Omega(g(n))

Summary sublinear (e.g. logarithmic), linear, and nlogn algorithms are good

polynomial algorithms may be acceptable in many cases exponential algorithms are bad

3 Algorithmic patterns

Recursion

recursion depth:

compilers/interpreters allocate space on a stack for the book-keeping for each function call

most environments limit the number of recursive calls: long chains of recursion are likely to cause errors

tail recursion:

e.g. linear search

easy to convert to iteration

easy to optimize, and optimized by many compilers (not by the Python interpreter)

Brute force in some cases we may need to enumerate all possible cases (e.g. to find the best solution)

common in combinatorial problems

often intractable, practical only for small input sizes

the beginning of finding a more efficient approach

example: segmentation

Divide and conquer divide the problem into smaller parts until it becomes trivial to solve

once small parts are solved, results are combined goes well with recursion

a particular flavor: binary search (sometimes called decrease and conquer)

example: nearest neighbors, merge sort, quick sort, integer multiplication, matrix multiplication, fast Fourier transform

not always yield good results, cost of merging should be less than the gain from the divisions

Greedy algorithm optimizes a local constraint results in correct solutions for some problems, in others they may result in "good enough" solutions

efficient if works

examples: graph algorithms (find shortest paths, scheduling)

e.g. produce minimum number of coins for a particular sum s (not correct for coins of 10, 30, 40 and sum value of 60)

Dynamic programming save earlier results to reduce computation

sometimes called memoization

examples: common parsing algorithms, Fibonacci

Others backtracking, branch-and-bound; randomized algorithms; distributed algorithms (sometimes called swarm optimization); transformation

4 Sorting

Bubble sort

compare first 2 elements, swap if not in order

shift and compare the next 2 elements, again swap if needed when read the end, repeat the process from the beginning unless there are no swaps in the last iteration

```
swapped = True
n = len(seq)
while swapped:
    swapped = False
    for i in range(n - 1):
        if seq[i] > seq[i + 1]:
            seq[i], seq[i + 1] = seq[i + 1], seq[i]
            swapped = True
```

concerns: many swaps, in-place

not practical, not used in practice

Insertion sort

assume the elements arrive one by one, and we have a sorted sequence

shift all elements larger than the new one to the right

```
put the new element in its correct place
for i in range(1, len(seq)):
    cur = seq[i]
    j = i
    while seq[j - 1] > cur:
        and j in range(1,i+1):
            seq[j] = seq[j - 1]
            j -= 1
    seq[j] = cur
```

performs reasonably fast for sorting short sequences, on longer sequences, performs worse than more advanced algorithms like merge sort or quick sort

in practice faster than bubble sort and selection sort

online: sort items as they arrive

stable: do not swap elements with equal keys

adaptive: faster if order of elements closer to sorted sequence

Merge sort

split the sequence

sort the subsequences

merge the sorted lists

```
# s1, s2: sequences to be merged
# s: target sequence
i, j = 0, 0
n = len(s1) + len(s2)
while i + j < n:
    if j == len(s2) or \
        i < len(s1) and s1[i] < s2[j]:
        s[i+j] = s1[i]
        i += 1
    else:
        s[i+j] = s2[j]
        j += 1
```

def merge_sort(s):

n = len(s)

if n <= 1: return

s1, s2 = s[:n//2], s[n//2:]

merge_sort(s1)

merge_sort(s2)

merge(s1, s2, s)

log n splits

particular useful for settings with low random-access memory, or sequential access

well-studied, many variants (in-place, non-recursive)

Quicksort

another popular divide-and-conquer sorting algorithm the big part of the work is done before splitting

worst time complexity is O(n²), but in practice performs better than merge sort on average

pick a pivot p, and divide the sequence into 3 parts:

L: smaller than p

G: larger than p

E: equal to p

sort L and G recursively

combination is simple concatenation

```
def qsort(seq):
    if len(seq <= 1): return seq
    return qsort([x for x in seq if x < seq[-1]]) + \
        [x for x in seq if x == seq[-1]] + \
        qsort([x for x in seq if x > seq[-1]])
```

similar to merge sort, performs O(n) operations at each level in recursion

overall complexity is proportional to n*depthOfTree

unlike merge sort, no balanced-tree guarantee, in the worst case the depth of the tree can be n, resulting in O(n²) complexity

worst case: when input sequence is sorted

randomized quicksort: pick the pivot randomly

best case: pick the median of the sequence as pivot, but finding median requires O(nlogn)

common approach: pick 3 values (typically first, middle, last) and select the median

can be easily implemented in-place

Algorithm	worst	average	best	memory	in-place	stable
Bubble sort	n ²	n ²	n	1	yes	yes
Insertion sort	n ²	n ²	n	1	yes	yes
Merge sort	n log n	n log n	n log n	n	no	yes
Quicksort	n log n	n log n	n log n	log n	yes	no

Bucket sort

puts elements of the input into a pre-defined number of

ordered 'buckets'

elements in each bucket is sorted (typically using insertion sort, worst case O(n²))

can retrieve the sorted elements by visiting each bucket

does not compare elements to each other when deciding which bucket to place them

in special cases results in O(n) worst-case complexity (as many buckets as keys)

Radix sort

sort objects with multiple keys

define the order of key pairs as (k1,1)<(k2,12) if k1<k2, or k1=k2 and 11<12, can be generalized to key tuples of any length

also known as lexicographic or dictionary order

use multiple stable bucket sorts for this purpose

5 Trees

a hierarchical, non-linear data structure

a graph with certain properties

in cl: parse trees, language trees, decision trees

Definitions

a set of nodes organized hierarchically with the following properties:

If a tree is non-empty, it has a special node root

Except the root node, every node in the tree has a unique parent (all nodes except the root are children of another node)

Alternatively, recursive definition:

the empty set of nodes is a tree

otherwise a tree contains a root with sub-trees as its children

siblings nodes with the same parent

internal nodes nodes with children

leaf nodes nodes without children

path sequence of connected nodes

ancestors/descendants any node in the path from the root to a particular node; a node is the descendant of its ancestors

internal nodes nodes with children

subtree a tree rooted by a non-root node

depth number of edges from root

height number of edges from the deepest descendant; the height of a tree is the height of its root

Ordered trees

if there is an ordering between siblings, e.g. document (e.g. HTML) structure tree, parse trees, family tree

in many cases order is not important (e.g. class hierarchy in object-oriented program, file tree in computer)

Binary trees

nodes can have at most 2 children

have a natural order: left/right child

proper/full: if every node has either 2 children or none complete: if every level except possibly the last is completely filled, all nodes at the last level is at the left

perfect: a full binary tree whose leaf nodes have the same depth

properties: for a binary tree with nl leaf nodes, ni internal nodes, n nodes and height h:

h+1 <= n <= 2^h(h+1)-1

1 <= nl <= 2^h

h <= ni <= 2^h-1

log(n+1)-1 <= h <= n-1

for any proper binary tree: nl=ni+1

Implementation

general case: linked data structure

arrays: root stored at index 0, left child of the node at index i stored at 2i+1, right child at 2i+2, parent at (i-1)/2 (if the binary tree is complete, this representation does not waste (much) space

Breadth first traversal (level order)

```
def breadth_first(root):
    queue = []
    queue.append(root)
    while queue:
        node = queue.pop(0)
        # process the node
        print(node.data)
        for child in node.children:
            queue.append(child)
```

Pre-order traversal

```
def pre_order(node):
    # process the node
    print(node.data)
    for child in node.children:
        pre_order(child)
```

Post-order traversal

```
def post_order(node):
    for child in node.children:
        post_order(child)
    # process the node
    print(node.data)
```

In-order traversal

```
def in_order(node):
    in_order(node.left)
    # process the node
    print(node.data)
    in_order(node.right)
```

6 Priority Queues

a collection, an abstract data type, that stores items (key-value pairs)

key: priority of the item

value: actual data of interest

its interface similar to a standard queue, but the item with the highest priority (minimum or maximum key value) instead of the first item entered into the queue is removed applications ranging from data compression to discrete optimization

Operations

insert(k,v): similar to enqueue(v)

remove(): similar to dequeue(), often called remove-min() or remove-max() depending on minimum or maximum key value that is considered having the highest priority

Implementation

unsorted list: insert O(1), remove O(n)

sorted list: insert O(1), remove O(1)

binary heap: insert O(logn), remove O(logn)

(some improvements:

d-ary heaps: insert O(log_d n), remove O(log_d n)

fibonacci heaps: insert O(1), remove O(logn)

Sorting with priority queues

implemented with sorted list = insertion sort O(n²)

use unsorted list = selection sort O(n²)

use binary heap = heap sort O(nlogn) (not stable; not in-place: needs O(n) extra space

in-place heap sort:

1. bottom-up heap construction

2. iteratively remove the maximum element and place it at the end

efficiency: heap construction O(n) + n*remove-min O(nlogn) = O(nlogn)

7 Binary Heaps

a binary tree where the nodes store items with an ordering relation

Properties

shape: a complete binary tree (all levels of the tree, except possibly the last one, are full; all empty slots (if any) are to the right of the filled nodes at the lowest level)

heap order: max-heap: parents' keys larger than children's keys; min-heap: parents' keys smaller than children's keys

Height log n

at least 2^h nodes -> h<=logn

at most 2^h(h+1)-1 nodes -> h>=log(n+1)-1

add new item

add to the first available slot

bubble up until the heap property is satisfied

at most h=logn comparisons/swaps

remove the min/max

item to be removed is at the root

ted graphs); complexity is $O(n+m)$ for n nodes and m edges
BFS

explore all options in parallel, divides the nodes into level
(starting node at level 0...)

typically implemented with a queue

if replace the queue with a stack, it is an iterative version of
DFS

```
def bfs(start):
    queue = [start]
    visited = {start: None}
    while queue:
        current = queue.pop(0)
        for node in current.neighbors():
            if node not in visited:
                visited[node] = current
                queue.append(node)
```

shortest path: if a node v is reachable from the start node, BFS finds the shortest path from the start node to v
complexity: $O(n+m)$

find path

traverse the graph from the source code, record the discover-
y edges

start from the target node, trace the path back to the source
with BFS, we get the shortest path
running time is the length of the path $O(n)$

```
def find_path(source, target, visited):
    path = []
    if target in visited:
        path.append(target)
        current = target
        while current is not source:
            parent = visited[current]
            path.append(parent)
            current = parent
        return path.reverse()
```

test connectivity, find connected com-
ponents, find cycle

connected: yes if the "visited"-nodes have the same length as
the graph nodes

find the connected components: run traversal multiple times
until all nodes are visited

cyclic: yes if there is a back edge during graph traversal

Directed graphs

terminology

for any pair of nodes u and v , a directed graph is

strongly connected: if there is a directed path between u to
 v and v to u

semi-connected: if there is a directed path between u to v or
 v to u

weakly connected: if the undirected graph obtained by replac-
ing all edges with undirected edges is a connected graph

check strong connectivity

naive approach: traverse the graph independently from each
node (strongly connected if all traversals visit all nodes)

time complexity $O(n*(n+m))$

better:

1. traverse the graph from an arbitrary node
2. reverse all edges, traverse again

3. intuition: if there a reverse path from D to A , then D is
reachable from A

time complexity: $O(n+m)$

transitive closure another graph where:

the set of nodes are the same as the original graph

there is an edge between two nodes u and v if v is reachable
from u

for undirected graph, can be computed by computing the
connected components

a straightforward algorithm:

1. run n graph traversals from each node in the graph
2. add an edge between the start node to any node discover-
ed by the traversal

time complexity $O(n*(n+m))$

(note: in a dense graph m is $O(n^2)$)

Floyd-Warshall algorithm: efficient if graph is implemented
with an adjacency matrix and is not sparse

1. setting transitive closure to the original graph
2. for $k=1,...,n$, add a directed edge (vi,vj) to transitive closure
if it already contains both (vi,vk) and (vk,vj)

```
T = [row[i] for row in G]
for k in range(n):
    for i in range(n):
        if i == k: continue
        for j in range(n):
            if j == i or j == k:
                continue
            T[i][j] = T[i][j] or \
                T[i][k] and T[k][j]
```

time complexity $O(n^3)$

a version of this is used for finding shortest paths in weigh-
ted graphs

Directed acyclic graphs(DAGs) directed graphs

without cycles

applications: course dependence, class inheritance, schedul-
ing constraints over tasks in a project, dependency parser
output

topological order: a sequence of nodes such that for every
directed edge (u,v) , u is listed before v ; there may be multi-
ple topological orderings, e.g. any acceptable order that the
courses can be taken

topological sort:

1. keep record of number of incoming edges
2. a node is ready to be placed in the sorted list if there are
no unprocessed incoming edges

time complexity $O(n+m)$

if topological ordering does not contain all the edges, the
graph includes a cycle

topo, ready = [], []

incount = {}

```
for u in nodes:
    incount[u] = u.undegree()
    if incount[u] == 0:
        ready.append(u)
while len(ready) > 0:
    u = ready.pop()
    topo.append(u)
    for v in u.neighbors():
        incount[v] -= 1
        if incount[v] == 0:
            ready.append(v)
```

Shortest Paths

weighted graph weights can be any numeric value, but
some algorithms require non-negative weights or
Euclidean weights (weights that are proper distance
metrics)

weights often indicate distance or cost, but can also
represent positive relations (e.g. affinity between nodes)

weight of a path: the sum of weights of the edges on the
path

shorted paths

applications: navigation, routing in computer networks,
optimal construction of electronic circuits, robotics,
transportation, finance...

shortest path on unweighted graphs: a BFS search tree

shortest path on unweighted graphs: different versions of
the problem, restrictions on weights

Dijkstra's algorithm a weighted version of BFS,
finds shorted path from a single source node to all
connected nodes

weights have to be non-negative

a greedy algorithm, grows a 'cloud' of nodes for which we
know the shortest paths from the source node
new nodes are included in the cloud in order of their
shortest paths from the source node

1. maintain a list D of minimum known distances to each
node
2. at each step:
-take closest node out of Q
-update the distances of all nodes
3. can be more efficient if Q is implemented using a priority
queue

```
1. D[s] ← 0
2. for each node v ≠ s do
    D[v] ← ∞
3. D[s] ← 0
4. Q ← nodes
5. while Q is not empty do
    6. Remove node u with min D[u] from Q
    7. for each edge (u, v) do
    8.     if D[u] + w(u, v) < D[v] then
    9.         D[v] ← D[u] + w(u, v)
10. D contains the shortest distances from s
```

complexity: in general $O(t_{find-min}n + t_{update-key}m)$

with list-based implementation of Q : $O((m+n)^2)=O(n^2)$

with a priority queue: $O((m+n)\log n)$
similar to traversal algorithms, does not give the shortest-
path tree, but can be extracted from distances D , running
time $O(n^2)$ or $O(n+m)$

1. $T \leftarrow \emptyset$
2. for $u \in D - \{s\}$ do
3. for each edge (v, u) do
4. if $D[v] > D[u] + w(v, u)$ then
5. $T \leftarrow T \cup (v, u)$

Shortest paths on DAGs directed acyclic graphs
similar to Dijkstra's, but simpler and faster
only difference is to follow a topological order
will also work with negative edge weights

Bellman-Ford algorithm single-source shortest
path problem for directed graph

include cycles, negative weights, exclude negative cycles
1. similar to earlier algorithms, initialize $D[s]=0$,
 $D[v]=\infty$

2. make n passes over the edges:
-update distances for each edge (relax edges)
-stop if there were no changes at the end of a pass

Minimum Spanning Tree
spanning tree

-spanning graph: includes all nodes

acyclic tree
minimum spanning tree

applications: network design, cluster analysis, traveling
salesman problem, object/network recognition in images,
avoiding cycles in broadcasting, dithering in images audio vid-
eo, error correction codes, dna sequencing

'cut property'

a cut of a graph is a partition that divides its nodes into two
disjoint (non-empty) sets

given any cut, the edge with the lowest weight across the cut
is in the MST

Prim-Jarnik algorithm greedy algorithm, finding
MST for weighted undirected graph

1. starts with a single 'start' node, grows the MST greedily

2. at each step:

-consider a cut between nodes visited and the rest of the no-
des

-select the minimum edge across the cut

3. repeat the process until all nodes are visited

1. pick any node s
2. $C[s] \leftarrow 0$
3. for each node $v \neq s$ do
4. $C[v] \leftarrow \infty$
5. $E[v] \leftarrow \text{None}$
6. $T \leftarrow \emptyset$
7. $Q \leftarrow \text{nodes}$
8. while Q is not empty do
9. Retrieve v with $\min C[v]$ from Q
10. Connect v to T
11. for edge (v, w) , where w is in Q do
12. if $\text{cost}(v, w) < C[w]$ then
13. $C[w] \leftarrow \text{cost}(v, w)$
14. $E[w] \leftarrow v$

complexity: two loops over number of nodes, $O(n^2)$ if we
need to search
if use a priority queue for Q , $O(m\log m)$
with a priority queue: $O((m+n)\log n)$
Kruskal's algorithm finding MST on undirected gra-
phs

1. start with each node in its own partition
 2. at each iteration, choose the edge with the minimum
weight across any two clusters, and join them
 3. terminates when there are no clusters to join
- ```
1. T ← ∅
2. for each node v do
 3. create_cluster(v)
4. for (u,v) in edges sorted by weight do
 5. if cluster(u) ≠ cluster(v) then
 6. T ← T ∪ {(u,v)}
 7. union(cluster(u), cluster(v))
```

loop over edges, but beware of the sorting requirement  
 $O(m\log m)$  with simple data structures

#### Directed trees

-rooted directed tree (arborescence): an acyclic directed  
graph where all nodes are reachable from the root node  
through a single directed path (what computational linguis-  
ticians simply calls a tree)

- anti-arborescence: a rooted directed tree where all edges  
are reversed

- polytree (a directed tree): a directed graph where undirec-  
ted edges form a tree

finding an MST in a directed graph = finding a rooted direc-  
ted tree

#### Chu-Liu/Edmonds algorithm

1. the MST for a directed graph has to start from a designa-  
ted root node
- if selected node has any incoming edges, remove them
- common practice to introduce an artificial root node with  
equal-weight edges to all nodes
2. for all non-root nodes, select the incoming edge with lo-  
west weight, remove others
3. if the resulting graph has no cycles, it is an MST 4. if there  
are cycles, break them
- consider the cycle as a single node - select the incoming edge  
that yields the lowest cost if used for breaking the cycle 5.  
repeat until no cycles remain

generally defined recursively: at each step, create a new  
graph with a contracted cycle call the procedure with the  
new graph

at most  $n$  recursions: the cycle has to include more nodes at  
every step

at each call,  $m$  steps for finding minimum incoming edge  
(also finding a cycle with  $O(n)$ , but  $m \gg n$ )  
the vanilla algorithm runs in  $O(mn)$ , there are improved ver-  
sions

in CL dependency parsing:

1. begin with fully connected weighted graph, except the  
root node has no incoming edges
2. weights are estimated from a treebank, typically determi-  
ned by a machine learning method trained on a treebank
3. often use probabilities, i.e. maximize the weight of the  
tree
4. one of the most common (and successful) approaches to  
dependency parsing

## 5 Maps & Hash Tables

### hash function

a one-way function that takes a variable-length object and  
turns it into a fixed-length bit string  
most common application: set, map (associative ar-  
ray/dictionary/symbol table)

other applic.: database indexing, cache management, effi-  
cient duplicate detection, file signature verification against  
corrupt/tempered files, password storage, electronic signa-  
tures, part of many cryptographic algorithms/applications

Set  
abstract data type, unordered collection without duplicates  
basic operations:  $x$  in  $s$ ,  $s.add(x)$ ,  $s.remove(x)$

map  
abstract data type, a collection that allows indexing with  
almost any data type (Python dict require immutable data  
type)

basic operations:  $d[key]$ ,  $d[key]=val$ ,  $d del d[key]$

implement sets and maps:

|                        | Check/retrieve | Add         | Remove      |
|------------------------|----------------|-------------|-------------|
| Sorted array:          | $O(\log n)$    | $O(n)$      | $O(n)$      |
| Unsorted array:        | $O(\log n)$    | $O(1)$      | $O(n)$      |
| Skip list:             | $O(\log n)$    | $O(\log n)$ | $O(\log n)$ |
| Balanced search trees: | $O(\log n)$    | $O(\log n)$ | $O(\log n)$ |
| Hash tables:           | $O(1)$         | $O(1)$      | $O(1)$      |

### hash function

$h()$  maps a key to an integer index between 0 and  $m$ (size of  
array)

we use  $h(k)$  as an index to an array(of size  $m$ )  
collision: occurs if 2 different key values are mapped to the  
same integer

2 parts: map any object (variable bit string) to an integer  
(e.g., 32 or 64 bit); compress the range of integers to map  
size( $m$ )

main challenge with implementing hash maps: avoid and  
handle the collisions

#### compress hash codes

easy way: use modulo  $m+1$  to map any integer to range  
 $[0,m]$

good hash functions minimize collisions, but collisions  
occur

2 common approaches to handle collisions: separate  
chaining, open addressing

separate chaining each array element keeps a pointer to  
a secondary container(typically a list); when an collision  
occurs, add the item to the list

complexity: all operations require locating the element first,  
cost include hashing(constant)+search in secondary data  
structure, worst-case  $O(n)$

with a good hash function, the probability of collision is  
 $n/m$ ,  $O(n/m)=O(1)$  (if  $m \gg n$ )

expected complexity for all operations is  $O(1)$

load factor = num of entries/num of indices  
low load factor: better run time (fewer collisions)/more  
memory usage

when load factor is over a threshold, the map is exten-  
ded(needs rehash)  
around 0.75 is considered optimal

open addressing(linear probing) during  
insertion, if there is a collision, look for the next empty slot  
and insert

during lookup, probe until there is an empty slot  
when delete an element, insert a special value that is treated  
as full during lookup and empty during insertion  
tends to create clusters of items, especially if load factor is  
high(>0.5)

quadratic probing provides some improvement  
quadratic probing prob  $(h(k)+i^2 \cdot m)$  mod  $m$  for  
 $i=0,1,...$  until an empty slot is found

if  $m$  is prime and load factor is less than 0.5, guaranteed to  
find an empty slot  
although better than linear probing, creates its own kind of  
clustering

double hashing prob  $(h(k)+i \cdot h'(k))$  mod  $m$  for  
 $i=0,1,...$ , where  $h'(k)$  is another hash function  
common choice:  $h'(k)=q-(k \bmod q)$  for a prime number  $q < m$

pseudo random number generator prob  
 $(h(k) + i \cdot r_i) \bmod m$  for  $i=0,1,...$ , where  $r_i$  is the  $i$ -th  
number generated by a pseudo number generator

pseudo random number generators generate numbers that  
are close to uniform, given the same seed, the sequence is  
deterministic

the most common choice for modern programming  
languages, avoids problems with inputs that intentionally  
generate hash collisions

#### has DoS attacks

a denial-of-service(DoS) attack aims to break or slow down  
an internet site/service  
input to a web-based program is passed as key-value pairs,  
which are typically stored in a dictionary

if one intentionally posss an input with large number of  
colliding keys, the hash table implementation needs to

chain long sequences or probe large number of times and  
eventually re-hash

this increases expected to  $O(1)$  time to worst-case complexi-  
ty

#### xor or add

hash codes must be consistent: if  $a=b$ ,  $h(a)=h(b)$   
should minimize collisions, values for  $h$  should be  
uniformly distributed  
should be fast to compute (or not if used for passwords  
simple approach: bitwise add (or XOR) each  $k$ -bit segment  
of the memory representation of the object, ignoring the  
overflow

in practice create many collisions due to associativity (abc,  
bca and cba get the same hash code)

#### polynomial hash codes

multiply with powers of a constant, will produce different  
values with sequences with the same items in different  
order

$$h(x) = \sum_{i=0}^n x_i a^{n-i-1} = x_0 a^{n-1} + x_1 a^{n-2} + \dots + x_{n-1}$$

#### cyclic-shift hash codes

shifts some bits from one end to the other at each step in  
the running sum

a fast way of obtaining a non-associative valid hash code  
since bitwise operations are simple

```
def cyclic_shift(s):
 mask = 0xffff
 h = 0
 for ch in s:
 h = (h << 5 & mask) | (h >> 11)
 h = ord(ch)
 return h
```

cryptographic hash functions in cryptography,  
it is important to have hash functions for which it is difficult  
to find two keys with the same hash value  
well-known hash functions: MD5, SHA-1, RIPEMD-160,  
Whirlpool, SHA-2, SHA-3, BLAKE2, BLAKE3  
designed for applications like digital fingerprinting,  
password storage  
computationally inefficient for use in data structures

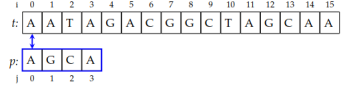
## 10 String Matching

find all occurrences of pattern  $p$  (length  $m$ ) in text  $t$  (length  
 $n$ )

The size of the alphabet ( $q$ ) is often an important factor  
 $p$  occurs in  $t$  with shift  $s$  if  $p[0:m] = t[s:s+m]$

A string  $x$  is a prefix/suffix of string  $y$ ; if  $y=xw$ /  $y=wx$  for a  
possibly empty string  $w$

#### Brute-force string search



Start from the beginning, of  $i = 0$  and  $j = 0$   
-if  $j == m$ , announce success with  $s = i$   
-if  $t[i] != p[j]$ : shift  $p$  (increase  $i$ , set  $j = 0$ )  
-otherwise: compare the next character (increase  $i$  and  $j$ ,  
repeat)

worst case: T: AAAAAAAAAAC, p: AAC  
Boyer-Moore algorithm  
start comparing from the end of  $p$

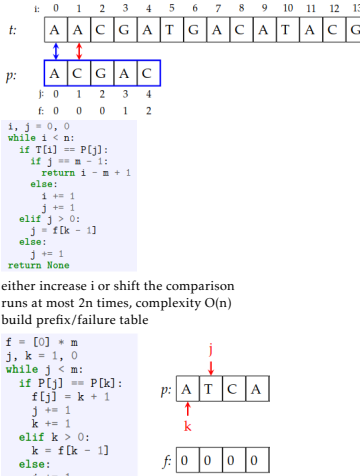
If  $t[i]$  does not occur in  $p$ , shift  $m$  steps

Otherwise, align the last occurrence of  $t[i]$  in  $p$  with  $t[i]$

```
last = {}
for j in range(m):
 last[p[j]] = j
i, j = m-1, m-1
while i < n:
 if t[i] == p[j]:
 if j == 0:
 return i
 else:
 i -= 1
 j -= 1
 else:
 k = last.get(t[i], -1)
 i += m - min(j, k+1)
 j = m-1
return None
```

on average performs better than brute-force  
fast case complexity  $O(nm)$ , e.g.  $t=aaaaaaaa$ ,  $p=baa$   
faster version exists  $O(n+m+q)$   
FSA

1. start at state 0, switch states based on the input
  2. all unspecified transitions go to state 0
  3. when at the accepting state, announce success
- naive attemp building automaton  $O(qm^3)$   
matching  $O(n)$   
space requirement  $O(qm)$  if stored in matrix  
faster algorithms for construction exists
- Knuth-Morris-Pratt (KMP) algorithm
1. in case of a match, increment both  $i$  and  $j$
  2. on failure, or at the end of the pattern, decide which new  
 $p[j]$  compare with  $t[i]$  based on a function  $f$
  3.  $f[j-1]$  tells which  $j$  value to resume the comparisons from



### Robin-Karp algorithm

1. instead of matching the string itself, matching the hash  
of it (based on a hash function)
2. If a match found, we need to verify – the match may be  
because of a hash collision
3. Otherwise, the algorithm makes a single comparison for  
each position in the text  
a hash should be computed for each position (size  $m$ )  
(rolling hash functions avoid this complication)  
rolling hash function changes the hash value only  
based on the item coming in and going out of the window  
to reduce collisions, better rolling-hash functions (e.g.  
polynomial hash functions) can be used

## 11 String Edit Distance

typically formulated as the (inverse) cost of obtaining one  
of the strings from the other through a number of edit  
operations

once we obtain the optimal edit operations, we may also be  
able to determine the optimal alignment between the  
strings

#### Hamming distance

number of different symbols in the corresponding positions  
easy calculation, but cannot handle sequences of different  
lengths

#### Longest common subsequence (LCS)

an order-preserving sequence of symbols from a string  
(solved by UNIX diff)

e.g. LCS(hygiene, hygeine) = hygine / hygiene

naive solution:

1. Enumerate all subsequences of the first string (exponential)
  2. Check if it is also a subsequence of the second string  
 $O(m^2 \cdot n)$  for strings of size  $n$  and  $m$
- recursive definition:
- $$LCS(Xx, Yy) = \begin{cases} LCS(X, Y)x & \text{if } x = y \\ \max(LCS(Xx, Y) \text{ and } LCS(X, Yy)) & \text{otherwise} \end{cases}$$

divide and conquer:

